

# NLP\_assignment1

April 10, 2025

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[3]: df = pd.read_csv("text_class - text_class.csv")
```

```
[39]: df
```

```
[39]:
```

	text	label
0	i loved the product, it's amazing!	positive
1	terrible service, i will never shop here again.	negative
2	the quality is good, but the delivery was late.	neutral
3	absolutely wonderful experience, highly recomm...	positive
4	product was damaged when it arrived, very disa...	negative
5	the customer support was very helpful and polite.	positive
6	worst purchase i've ever made.	negative
7	satisfied with the product but the price was t...	neutral

```
[11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8 entries, 0 to 7
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    text    8 non-null         object
1    label    8 non-null         object
dtypes: object(2)
memory usage: 260.0+ bytes
```

```
[17]: # Print the total number of rows
print("Total number of rows:", len(df))

# Print the count of unique labels
print("Count of unique labels:", df[['text', 'label']].nunique())
```

```
Total number of rows: 8
Count of unique labels: text      8
```

```
label      3
dtype: int64
```

```
[29]: columns_to_lower = ['text', 'label']
      for col in columns_to_lower:
          df[col] = df[col].str.lower()
```

```
[31]: df.head()
```

```
[31]:
```

	text	label
0	i loved the product, it's amazing!	positive
1	terrible service, i will never shop here again.	negative
2	the quality is good, but the delivery was late.	neutral
3	absolutely wonderful experience, highly recomm...	positive
4	product was damaged when it arrived, very disa...	negative

```
[3]: import pandas as pd
      import string
      import nltk
      from nltk.corpus import stopwords
      from nltk.tokenize import word_tokenize

      nltk.download('punkt')
      nltk.download('stopwords')

      def preprocess_text(text):
          # Convert text to lowercase
          text = text.lower()
          # Remove punctuation and special characters
          text = text.translate(str.maketrans('', '', string.punctuation))
          # Tokenize the text
          tokens = word_tokenize(text)
          # Remove stopwords
          stop_words = set(stopwords.words('english'))
          filtered_tokens = [word for word in tokens if word not in stop_words]
          return ' '.join(filtered_tokens) # Join tokens back into a string

      data = pd.read_csv("text_class - text_class.csv")

      # Apply preprocessing to each text column
      data['processed_text'] = data['text'].apply(preprocess_text)

      # Display the processed dataset
      print(data)
```

[nltk\_data] Downloading package punkt to

```
[nltk_data] C:\Users\numan\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\numan\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

	text	label
0	I loved the product, it's amazing!	positive
1	Terrible service, I will never shop here again.	negative
2	The quality is good, but the delivery was late.	neutral
3	Absolutely wonderful experience, highly recomm...	positive
4	Product was damaged when it arrived, very disa...	negative
5	The customer support was very helpful and polite.	positive
6	Worst purchase I've ever made.	negative
7	Satisfied with the product but the price was t...	neutral

	processed_text
0	loved product amazing
1	terrible service never shop
2	quality good delivery late
3	absolutely wonderful experience highly recommend
4	product damaged arrived disappointed
5	customer support helpful polite
6	worst purchase ive ever made
7	satisfied product price high

```
[7]: from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Splitting data
X = data['processed_text']
y = data['label']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    random_state=42)

# Convert text data to numerical representation (Bag-of-Words)
vectorizer = CountVectorizer()
X_train_vect = vectorizer.fit_transform(X_train)
X_test_vect = vectorizer.transform(X_test)

# Training model
model = LogisticRegression()
model.fit(X_train_vect, y_train)

# Prediction on test data
y_pred = model.predict(X_test_vect)
```

```
# Accuracy calculation and output
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")
```

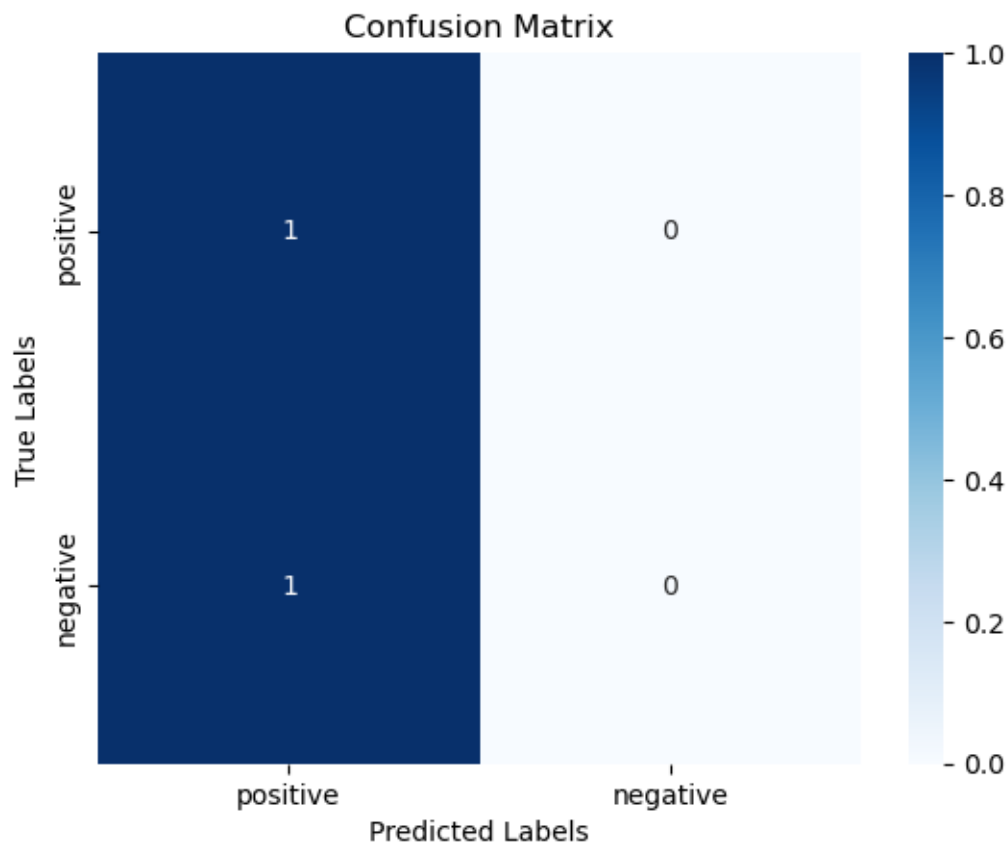
Accuracy: 50.00%

```
[11]: from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
import matplotlib.pyplot as plt

# confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred, labels=['positive', 'negative'])

# Display the heatmap
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['positive', 'negative'], yticklabels=['positive', 'negative'])
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()

# classification report for performance metrics
print("Classification Report:")
print(classification_report(y_test, y_pred, target_names=['positive',
            'negative']))
```



Classification Report:

	precision	recall	f1-score	support
positive	0.00	0.00	0.00	1
negative	0.50	1.00	0.67	1
accuracy			0.50	2
macro avg	0.25	0.50	0.33	2
weighted avg	0.25	0.50	0.33	2

C:\Users\numan\anaconda3\Lib\site-packages\sklearn\metrics\\_classification.py:1509: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))

C:\Users\numan\anaconda3\Lib\site-packages\sklearn\metrics\\_classification.py:1509: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```
C:\Users\numan\anaconda3\Lib\site-  
packages\sklearn\metrics\_classification.py:1509: UndefinedMetricWarning:  
Precision is ill-defined and being set to 0.0 in labels with no predicted  
samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

The confusion matrix provides detailed insights into the model's performance by showing the number of correct and incorrect predictions for each class. It helps identify specific areas where the model struggles, such as frequent misclassification between similar classes, enabling targeted improvements in data preprocessing or model refinement

[ ]: